

Single studies using the CohortMethod package

Martijn J. Schuemie, Marc A. Suchard and Patrick Ryan

2015-03-27

Contents

1	Introduction	1
2	Installation instructions	2
3	Data extraction	2
3.1	Configuring the connection to the server	2
3.2	Preparing the exposures and outcome(s)	2
3.3	Extracting the data from the server	4
4	Propensity scores	7
4.1	Fitting a propensity model	7
4.2	Propensity score diagnostics	7
4.3	Using the propensity score	8
4.4	Evaluating covariate balance	11
5	Outcome models	13
5.1	Fitting the outcome model	13
5.2	Inspecting the outcome model	14
5.3	Kaplan-Meier plot	15
5.4	Attrition diagram	16
6	Acknowledgments	16

1 Introduction

This vignette describes how you can use the `CohortMethod` package to perform a single new-user cohort study. We will walk through all the steps needed to perform an exemplar study, and we have selected the well-studied topic of the effect of coxibs versus non-selective nonsteroidal anti-inflammatory drugs (NSAIDs) on gastrointestinal (GI) bleeding-related hospitalization. For simplicity, we focus on one coxib – celecoxib – and one non-selective NSAID – diclofenac.

2 Installation instructions

Before installing the `CohortMethod` package make sure you have Java available. Java can be downloaded from www.java.com. For Windows users, RTools is also necessary. RTools can be downloaded from [CRAN](http://cran.r-project.org/web/packages/Rtools/index.html).

The `CohortMethod` package is currently maintained in a [Github repository](#), and has dependencies on other packages in Github. All of these packages can be downloaded and installed from within R using the `devtools` package:

```
install.packages("devtools")
library(devtools)
install_github("ohdsi/SqlRender")
install_github("ohdsi/DatabaseConnector")
install_github("ohdsi/Cyclops")
install_github("ohdsi/CohortMethod")
```

Once installed, you can type `library(CohortMethod)` to load the package.

3 Data extraction

The first step in running the `CohortMethod` is extracting all necessary data from the database server holding the data in the Observational Medical Outcomes Partnership (OMOP) Common Data Model (CDM) format.

3.1 Configuring the connection to the server

We need to tell R how to connect to the server where the data are. `CohortMethod` uses the `DatabaseConnector` package, which provides the `createConnectionDetails` function. Type `?createConnectionDetails` for the specific settings required for the various database management systems (DBMS). For example, one might connect to a PostgreSQL database using this code:

```
connectionDetails <- createConnectionDetails(dbms = "postgresql",
                                             server = "localhost/ohdsi",
                                             user = "joe",
                                             password = "supersecret")

cdmDatabaseSchema <- "my_cdm_data"
resultsDatabaseSchema <- "my_results"
```

The last two lines define the `cdmDatabaseSchema` and `resultSchema` variables, which we'll use later to tell R where the data in CDM format live, and where we want to write intermediate and result tables. Note that for Microsoft SQL Server, these variables need to specify both the database and the schema, so for example `cdmDatabaseSchema <- "my_cdm_data.dbo"`.

3.2 Preparing the exposures and outcome(s)

We need to define the exposures and outcomes for our study. We do this by writing SQL statements against the OMOP CDM that populate a table of events in which we are interested. For our example study, we have created a file called *coxibVsNonselVsGiBleed.sql* with the following contents:

```

/*****
File coxibVsNonselVsGiBleed.sql
*****/

IF OBJECT_ID('@resultsDatabaseSchema.coxibVsNonselVsGiBleed', 'U') IS NOT NULL
    DROP TABLE @resultsDatabaseSchema.coxibVsNonselVsGiBleed;

CREATE TABLE @resultsDatabaseSchema.coxibVsNonselVsGiBleed (
    cohort_definition_id INT,
    cohort_start_date DATE,
    cohort_end_date DATE,
    subject_id BIGINT
);

INSERT INTO @resultsDatabaseSchema.coxibVsNonselVsGiBleed (
    cohort_definition_id,
    cohort_start_date,
    cohort_end_date,
    subject_id
)
SELECT 1, -- Exposure
    drug_era_start_date,
    drug_era_end_date,
    person_id
FROM @cdmDatabaseSchema.drug_era
WHERE drug_concept_id = 1118084; -- celecoxib

INSERT INTO @resultsDatabaseSchema.coxibVsNonselVsGiBleed (
    cohort_definition_id,
    cohort_start_date,
    cohort_end_date,
    subject_id
)
SELECT 2, -- Comparator
    drug_era_start_date,
    drug_era_end_date,
    person_id
FROM @cdmDatabaseSchema.drug_era
WHERE drug_concept_id = 1124300; -- diclofenac

INSERT INTO @resultsDatabaseSchema.coxibVsNonselVsGiBleed (
    cohort_definition_id,
    cohort_start_date,
    cohort_end_date,
    subject_id
)
SELECT 3, -- Outcome
    condition_start_date,
    condition_end_date,
    condition_occurrence.person_id
FROM @cdmDatabaseSchema.condition_occurrence
INNER JOIN @cdmDatabaseSchema.visit_occurrence
    ON condition_occurrence.visit_occurrence_id = visit_occurrence.visit_occurrence_id

```

```
WHERE condition_concept_id IN (
  SELECT descendant_concept_id
  FROM @cdmDatabaseSchema.concept_ancestor
  WHERE ancestor_concept_id = 192671 -- GI - Gastrointestinal haemorrhage
)
AND visit_occurrence.place_of_service_concept_id IN (9201, 9203);
```

This is parameterized SQL which can be used by the `SqlRender` package. We use parameterized SQL so we do not have to pre-specify the names of the CDM and result schemas. That way, if we want to run the SQL on a different schema, we only need to change the parameter values; we do not have to change the SQL code. By also making use of translation functionality in `SqlRender`, we can make sure the SQL code can be run in many different environments.

```
library(SqlRender)
sql <- readSql("coxibVsNonselVsGiBleed.sql")
sql <- renderSql(sql,
  cdmDatabaseSchema = cdmDatabaseSchema,
  resultsDatabaseSchema = resultsDatabaseSchema)$sql
sql <- translateSql(sql, targetDialect = connectionDetails$dbms)$sql

connection <- connect(connectionDetails)
executeSql(connection, sql)
```

In this code, we first read the SQL from the file into memory. In the next line, we replace the two parameter names with the actual values. We then translate the SQL into the dialect appropriate for the DBMS we already specified in the `connectionDetails`. Next, we connect to the server, and submit the rendered and translated SQL.

If all went well, we now have a table with the events of interest. We can see how many events per type:

```
sql <- paste("SELECT cohort_definition_id, COUNT(*) AS count",
  "FROM @resultsDatabaseSchema.coxibVsNonselVsGiBleed",
  "GROUP BY cohort_definition_id")
sql <- renderSql(sql, resultsDatabaseSchema = resultsDatabaseSchema)$sql
sql <- translateSql(sql, targetDialect = connectionDetails$dbms)$sql

querySql(connection, sql)
```

```
#> cohort_concept_id count
#> 1                1 128785
#> 2                2 417027
#> 3                3 419141
```

3.3 Extracting the data from the server

Now we can tell `CohortMethod` to define the cohorts based on our events, and extract all necessary data for our analysis:

```
# Get all NSAID Concept IDs for exclusion:
sql <- paste("SELECT concept_id",
  "FROM @cdmDatabaseSchema.concept_ancestor",
  "INNER JOIN @cdmDatabaseSchema.concept",
```

```

      "ON descendant_concept_id = concept_id",
      "WHERE ancestor_concept_id = 21603933")
sql <- SqlRender::renderSql(sql, cdmDatabaseSchema = cdmDatabaseSchema)$sql
sql <- SqlRender::translateSql(sql, targetDialect = connectionDetails$dbms)$sql
nsaids <- querySql(connection, sql)
nsaids <- nsaids$CONCEPT_ID

```

```

#Load data:
cohortData <- getDbCohortData(connectionDetails,
                             cdmDatabaseSchema = cdmDatabaseSchema,
                             oracleTempSchema = resultsDatabaseSchema,
                             targetDrugConceptId = 1,
                             comparatorDrugConceptId = 2,
                             indicationConceptIds = c(),
                             washoutWindow = 183,
                             indicationLookbackWindow = 183,
                             studyStartDate = "",
                             studyEndDate = "",
                             exclusionConceptIds = nsaids,
                             outcomeConceptIds = 3,
                             outcomeConditionTypeConceptIds = c(),
                             exposureDatabaseSchema = resultsDatabaseSchema,
                             exposureTable = "coxibVsNonseIvsGiBleed",
                             outcomeDatabaseSchema = resultsDatabaseSchema,
                             outcomeTable = "coxibVsNonseIvsGiBleed",
                             useCovariateDemographics = TRUE,
                             useCovariateConditionOccurrence = TRUE,
                             useCovariateConditionOccurrence365d = TRUE,
                             useCovariateConditionOccurrence30d = TRUE,
                             useCovariateConditionOccurrenceInpt180d = TRUE,
                             useCovariateConditionEra = TRUE,
                             useCovariateConditionEraEver = TRUE,
                             useCovariateConditionEraOverlap = TRUE,
                             useCovariateConditionGroup = TRUE,
                             useCovariateDrugExposure = TRUE,
                             useCovariateDrugExposure365d = TRUE,
                             useCovariateDrugExposure30d = TRUE,
                             useCovariateDrugEra = TRUE,
                             useCovariateDrugEra365d = TRUE,
                             useCovariateDrugEra30d = TRUE,
                             useCovariateDrugEraEver = TRUE,
                             useCovariateDrugEraOverlap = TRUE,
                             useCovariateDrugGroup = TRUE,
                             useCovariateProcedureOccurrence = TRUE,
                             useCovariateProcedureOccurrence365d = TRUE,
                             useCovariateProcedureOccurrence30d = TRUE,
                             useCovariateProcedureGroup = TRUE,
                             useCovariateObservation = TRUE,
                             useCovariateObservation365d = TRUE,
                             useCovariateObservation30d = TRUE,
                             useCovariateObservationBelow = TRUE,
                             useCovariateObservationAbove = TRUE,
                             useCovariateObservationCount365d = TRUE,

```

```

        useCovariateConceptCounts = TRUE,
        useCovariateRiskScores = TRUE,
        useCovariateInteractionYear = FALSE,
        useCovariateInteractionMonth = FALSE,
        excludedCovariateConceptIds = nsaid,
        excludeDrugsFromCovariates = FALSE,
        deleteCovariatesSmallCount = 100)

cohortData

```

```

#> CohortData object
#>
#> Treatment concept ID: 1
#> Comparator concept ID: 2
#> Outcome concept ID(s): 3

```

There are many parameters, but they are all documented in the `CohortMethod` manual. In short, we are pointing the function to the table created earlier and indicating which concept IDs in that table identify the target, comparator and outcome. Note that in this example, we do not restrict the study to people having a particular indication via `indicationConceptIds = c()`, but this is something you would often want to do. We do instruct that people with prior exposure to any NSAID should be excluded, and that many different covariates should be constructed, including covariates for all conditions, drug exposures, and procedures that were found on or before the index date.

All data about the cohorts, outcomes, and covariates are extracted from the server and stored in the `cohortData` object. This object uses the package `ff` to store information in a way that ensures R does not run out of memory, even when the data are large.

We can use the generic `summary()` function to view some more information of the data we extracted:

```
summary(cohortData)
```

```

#> CohortData object summary
#>
#> Treatment concept ID: 1
#> Comparator concept ID: 2
#> Outcome concept ID(s): 3
#>
#> Treated persons: 13207
#> Comparator persons: 59619
#>
#> Outcome counts:
#>   Event count Person count
#> 3         2785         1884
#>
#> Covariates:
#> Number of covariates: 24049
#> Number of non-zero covariate values: 47871135

```

3.3.1 Saving the data to file

Creating the `cohortData` file can take considerable computing time, and it is probably a good idea to save it for future sessions. Because `cohortData` uses `ff`, we cannot use R's regular save function. Instead, we'll have to use the `saveCohortData()` function:

```
saveCohortData(cohortData, "coxibVsNonselVsGiBleed")
```

We can use the `loadCohortData()` function to load the data in a future session.

4 Propensity scores

The `CohortMethod` can use propensity scores to adjust for potential confounders. Instead of the traditional approach of using a handfull of predefined covariates, `CohortMethod` typically uses thousands to millions of covariates that are automatically constructed based on conditions, procedures and drugs in the records of the subjects.

4.1 Fitting a propensity model

We can fit a propensity model using the covariates constructed by the `getDbCohortData()` function:

```
ps <- createPs(cohortData, outcomeConceptId = 3)
```

The `createPs()` function uses the `Cyclops` package to fit a large-scale regularized logistic regression. Note that we have to tell `createPs()` what the `outcomeConceptId` is for which we will use the model so it can remove subjects who had the outcome prior to the index date before fitting the model.

To fit the propensity model, `Cyclops` needs to know the hyperparameter value which specifies the variance of the prior. By default `Cyclops` will use cross-validation to estimate the optimal hyperparameter. However, be aware that this can take a really long time. You can use the `prior` and `control` parameters of the `createPs()` to specify `Cyclops` behaviour, including using multiple CPUs to speed-up the cross-validation.

4.2 Propensity score diagnostics


We can compute the area under the receiver-operator curve (AUC) for the propensity score model:

```
computePsAuc(ps)
```

```
#> [1] 0.8563299
```

We can also plot the propensity score distribution, although we prefer the preference score distribution:

```
plotPs(ps, scale = "preference")
```



It is also possible to inspect the propensity model itself by showing the covariates that have non-zero coefficients:

```
propensityModel <- getPsModel(ps, cohortData)
head(propensityModel)
```


```
#>      coefficient      id      covariateName
#> 1312  -3.2440631 1150871503 ... with cohort index: 1150871-Misoprostol
#> 1      -1.0783597      13      Age group: 15-19
#> 101    -1.0434257      12      Age group: 10-14
#> 328     0.9427427      2007     Index year: 2007
#> 1371  -0.8636176 19102773402 ...in 5 MG/ML Ophthalmic Solution [Vigamox]
#> 1558   0.8340975 40664232701 ...IC TREATMENT VISIT (AS PART OF CONTRACT)
```

One advantage of using the regularization when fitting the propensity model is that most coefficients will shrink to zero and fall out of the model. It is a good idea to inspect the remaining variables for anything that should not be there, for example instrumental variables.

4.3 Using the propensity score


We can use the propensity scores to trim, stratify, or match our population. For example, one could trim to equipoise, meaning only subjects with a preference score between 0.25 and 0.75 are kept:


```
psTrimmed <- trimByPsToEquipoise(ps)
plotPs(psTrimmed, ps, scale = "preference")
```




Instead (or additionally), we could stratify the population based on the propensity score:

```
psStratified <- stratifyByPs(ps, numberOfStrata = 5)
plotPs(psStratified, ps, scale = "preference")
```



We can also match subjects based on propensity scores. In this example, we're using one-to-one matching:

```
strata <- matchOnPs(ps, caliper = 0.25, caliperScale = "standardized", maxRatio = 1)
plotPs(strata, ps)
```




Note that for both stratification and matching it is possible to specify additional matching criteria such as age and sex using the `stratifyByPsAndCovariates()` and `matchOnPsAndCovariates()` functions, respectively.

4.4 Evaluating covariate balance

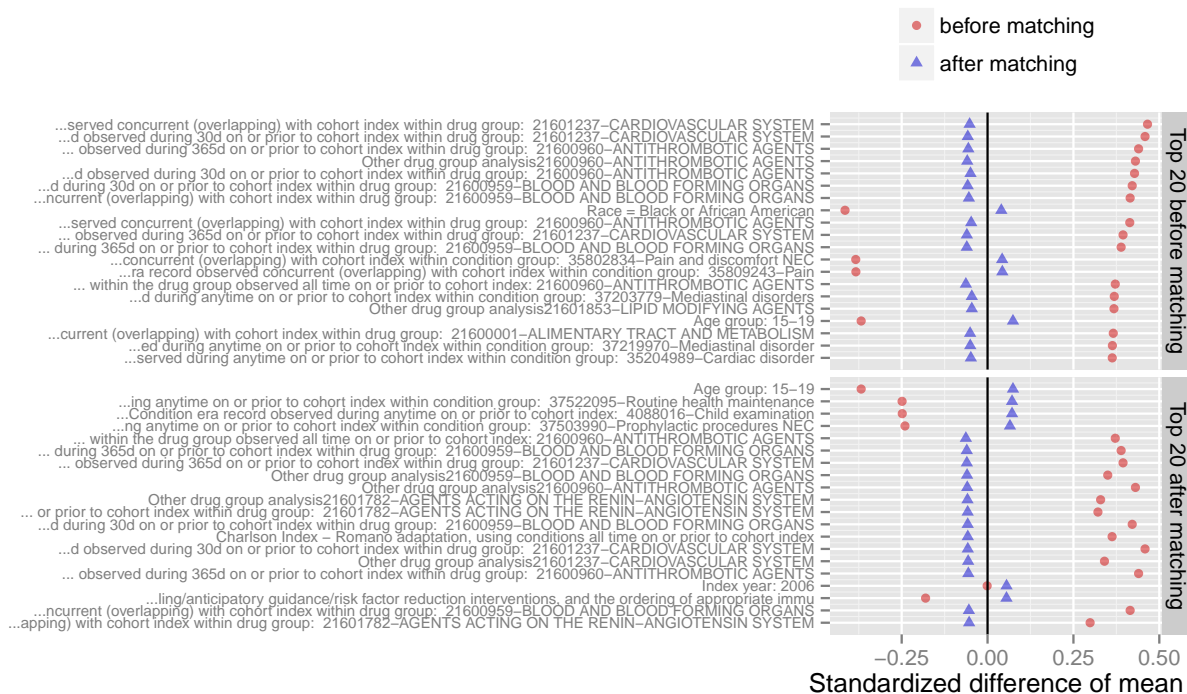
To evaluate whether our use of the propensity score is indeed making the two cohorts more comparable, we can compute the covariate balance before and after trimming, matching, and/or stratifying:

```
balance <- computeCovariateBalance(strata, cohortData, outcomeConceptId = 3)
```

```
plotCovariateBalanceScatterPlot(balance)
```



```
plotCovariateBalanceOfTopVariables(balance)
```



5 Outcome models

The outcome model is a model describing which variables are associated with the outcome.

5.1 Fitting the outcome model

In theory we could fit an outcome model without using the propensity scores. In this example we are fitting an outcome model using a Cox regression. The risk window is defined as time of exposure + 30 days:

```
outcomeModel <- fitOutcomeModel(outcomeConceptId = 3,
                                cohortData = cohortData,
                                riskWindowStart = 0,
                                riskWindowEnd = 30,
                                addExposureDaysToEnd = TRUE,
                                useCovariates = FALSE,
                                modelType = "cox",
                                stratifiedCox = FALSE)

outcomeModel
```

```
#> Model type: cox
#> Status: OK
#>
#> Prior variance: Inf
#>           Estimate lower .95 upper .95      logRr seLogRr
#> treatment  0.939443  0.713658  1.227566 -0.062468  0.1365
```

But of course we want to make use of the matching done on the propensity score:

```
outcomeModel <- fitOutcomeModel(outcomeConceptId = 3,
                                cohortData = cohortData,
                                subPopulation = strata,
                                riskWindowStart = 0,
                                riskWindowEnd = 30,
                                addExposureDaysToEnd = TRUE,
                                useCovariates = FALSE,
                                modelType = "cox",
                                stratifiedCox = TRUE)

outcomeModel
```

```
#> Model type: cox
#> Status: OK
#>
#> Prior variance: Inf
#>           Estimate lower .95 upper .95      logRr seLogRr
#> treatment  0.60465   0.36705   0.97675 -0.50310  0.2447
```

Note that we define the subpopulation to be only those in the `strata` object, which we created earlier by matching on the propensity score. We also now use a stratified Cox model, conditioning on the propensity score match sets.

One final refinement would be to use the same covariates we used to fit the propensity model to also fit the outcome model. This way we are more robust against misspecification of the model, and more likely to

remove bias. For this we use the regularized Cox regression in the `Cyclops` package. (Note that the treatment variable is automatically excluded from regularization.)

```
outcomeModel <- fitOutcomeModel(outcomeConceptId = 3,
                                cohortData = cohortData,
                                subPopulation = strata,
                                riskWindowStart = 0,
                                riskWindowEnd = 30,
                                addExposureDaysToEnd = TRUE,
                                useCovariates = TRUE,
                                modelType = "cox",
                                stratifiedCox = TRUE)

outcomeModel
```

```
#> Model type: cox
#> Status: OK
#>
#> Prior variance: 0.0416455739602391
#>           Estimate lower .95 upper .95    logRr seLogRr
#> treatment  0.79439    0.42584    1.47257 -0.23018  0.3149
```

5.2 Inspecting the outcome model

We can inspect more details of the outcome model:

```
summary(outcomeModel)
```

```
#> Model type: cox
#> Status: OK
#>
#> Counts
#>           Comparator Treated
#> Nr. of persons      10714   10713
#> Nr. of events         62     64
#> Person time (days)  915975 1418302
#>
#> Model
#>           Nr. of betas Nr. of non-zero betas      Number of strata
#>           18862              10              10724
#>
#> Coefficients
#>           Estimate lower .95 upper .95    logRr seLogRr
#> treatment  0.79439    0.42584    1.47257 -0.23018  0.3149
#>
#> Prior variance: 0.0416455739602391
```

```
coef(outcomeModel)
```

```
#> [1] -0.2301838
```

```
confint(outcomeModel)
```

```
#> [1] -0.8536895  0.3870119
```

We can also see the covariates that ended up in the outcome model:

```
fullOutcomeModel <- getOutcomeModel(outcomeModel, cohortData)
head(fullOutcomeModel)
```


```
#>   coefficient      id covariateName
#> 5 -0.31699629 21602796601 ...21602796-ANTIBACTERIALS FOR SYSTEMIC USE
#> 8  0.24262483      1100 ...ons all time on or prior to cohort index
#> 7 -0.23018381        1 Cohort definition ID
#> 9  0.14379871 21602098601 ...ICOSTEROIDS, DERMATOLOGICAL PREPARATIONS
#> 10 0.09652774      1101 ...ons all time on or prior to cohort index
#> 2  0.05851533      1001 ...rved in 365d on or prior to cohort index
```

5.3 Kaplan-Meier plot

We can create the Kaplan-Meier plot:

```
plotKaplanMeier(outcomeModel, includeZero = FALSE)
```

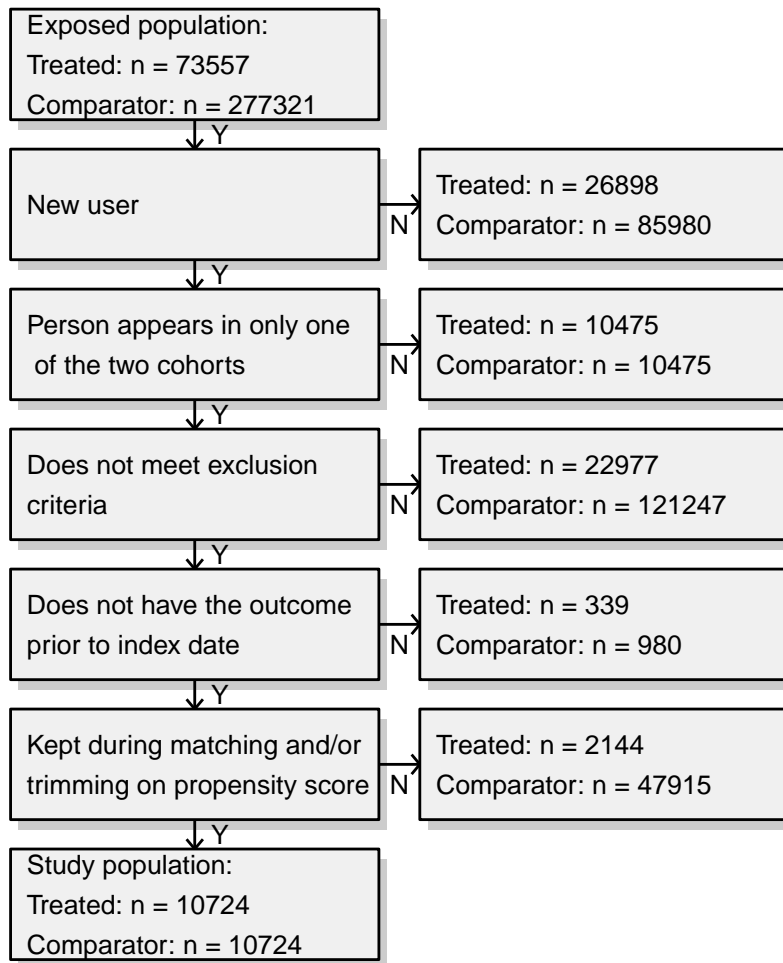
```
#> Warning in plotKaplanMeier(outcomeModel, includeZero = FALSE): The outcome
#> model is stratified, but the stratification is not visible in the plot
```



5.4 Attrition diagram

We can also investigate how we got to the study population by drawing the attrition diagram:

```
drawAttritionDiagram(outcomeModel)
```



6 Acknowledgments

Considerable work has been dedicated to provide the `CohortMethod` package.

```
citation("CohortMethod")
```

```
#>  
#> To cite package 'CohortMethod' in publications use:
```



```

#>
#> Martijn J. Schuemie, Marc A. Suchard and Patrick B. Ryan (2015).
#> CohortMethod: New-user cohort method with large scale propensity
#> and outcome models. R package version 1.0.3.
#>
#> A BibTeX entry for LaTeX users is
#>
#> @Manual{,
#>   title = {CohortMethod: New-user cohort method with large scale propensity and outcome models},
#>   author = {Martijn J. Schuemie and Marc A. Suchard and Patrick B. Ryan},
#>   year = {2015},
#>   note = {R package version 1.0.3},
#> }
#>
#> ATTENTION: This citation information has been auto-generated from
#> the package DESCRIPTION file and may need manual editing, see
#> 'help("citation")'.

```

Further, CohortMethod makes extensive use of the Cyclops package.

```
citation("Cyclops")
```

```

#>
#> To cite Cyclops in publications use:
#>
#> Suchard MA, Simpson SE, Zorych I, Ryan P and Madigan D (2013).
#> "Massive parallelization of serial inference algorithms for
#> complex generalized linear models." _ACM Transactions on Modeling
#> and Computer Simulation_, *23*, pp. 10. <URL:
#> http://dl.acm.org/citation.cfm?id=2414791>.
#>
#> A BibTeX entry for LaTeX users is
#>
#> @Article{,
#>   author = {M. A. Suchard and S. E. Simpson and I. Zorych and P. Ryan and D. Madigan},
#>   title = {Massive parallelization of serial inference algorithms for complex generalized linear m
#>   journal = {ACM Transactions on Modeling and Computer Simulation},
#>   volume = {23},
#>   pages = {10},
#>   year = {2013},
#>   url = {http://dl.acm.org/citation.cfm?id=2414791},
#> }

```

This work is supported in part through the National Science Foundation grant IIS 1251151.